

## K-Neighbors

```
In [1]: import numpy as np
import pandas as pd
```

### Dataset Description

#### Electrical Grid Stability Simulated Data Set

The local stability analysis of the 4-node star system (electricity producer in the center) implementing Decentralised Smart Grid Control concept.

```
In [2]: data=pd.read_csv("data.csv")
```

```
In [3]: data.head()
```

```
Out[3]:
```

	tau1	tau2	tau3	tau4	p1	p2	p3	p4	g1	g2	g3	
0	2.959060	3.079885	8.381025	9.780754	3.763085	-0.782604	-1.257395	-1.723086	0.650456	0.859578	0.887445	0.9580
1	9.304097	4.902524	3.047541	1.369357	5.067812	-1.940058	-1.872742	-1.255012	0.413441	0.862414	0.562139	0.7817
2	8.971707	8.848428	3.046479	1.214518	3.405158	-1.207456	-1.277210	-0.920492	0.163041	0.766689	0.839444	0.1098
3	0.716415	7.669600	4.486641	2.340563	3.963791	-1.027473	-1.938944	-0.997374	0.446209	0.976744	0.929381	0.3627
4	3.134112	7.608772	4.943759	9.857573	3.525811	-1.125531	-1.845975	-0.554305	0.797110	0.455450	0.656947	0.8209

```
In [4]: from sklearn.preprocessing import StandardScaler,LabelBinarizer
from sklearn.model_selection import train_test_split
lb=LabelBinarizer()
sc=StandardScaler()
```

```
In [5]: X=data.iloc[:, :-1]
Y=data.iloc[:, -1]
```

```
In [6]: X=sc.fit_transform(X)
Y=lb.fit_transform(Y)
```

```
In [7]: X_train,X_test,Y_train,Y_test=train_test_split(X,Y,test_size=0.2,random_state=0)
```

### K-Neighbours (Library)

```
In [8]: from sklearn.neighbors import KNeighborsClassifier
from sklearn.metrics import classification_report
knn=KNeighborsClassifier(n_neighbors=3,metric='euclidean')
```

```
In [9]: knn.fit(X_train,Y_train)
```

C:\Anaconda\lib\site-packages\ipykernel\_launcher.py:1: DataConversionWarning: A column-vector y was passed when a 1d array was expected. Please change the shape of y to (n\_samples, ), for example using ravel().  
 """Entry point for launching an IPython kernel.

```
Out[9]: KNeighborsClassifier(algorithm='auto', leaf_size=30, metric='euclidean',
metric_params=None, n_jobs=None, n_neighbors=3, p=2,
weights='uniform')
```

```
In [10]: Y_pred=knn.predict(X_test)
```

```
In [11]: print(classification_report(Y_test,Y_pred))
print("Accuracy: {0:.2f} %".format(knn.score(X_test,Y_test)*100))
```

	precision	recall	f1-score	support
0	0.91	0.86	0.88	727
1	0.92	0.95	0.94	1273
accuracy			0.92	2000
macro avg	0.91	0.90	0.91	2000
weighted avg	0.92	0.92	0.92	2000

Accuracy: 91.65 %

## K-Neighbours (Custom)

```
In [12]: class KNeighbours(object):
def __init__(self, k):
    self.k = k

    @staticmethod
    def euclid_dist(v1, v2):
        v1, v2 = np.array(v1), np.array(v2)
        distance = 0
        for i in range(len(v1) - 1):
            distance += (v1[i] - v2[i]) ** 2
        return np.sqrt(distance)

    def predict(self, train_set, test_inst):
        distances = []
        for i in range(len(train_set)):
            dist = self.euclid_dist(train_set[i][::-1], test_inst)
            distances.append((train_set[i], dist))
        distances.sort(key=lambda x: x[1])

        neighbours = []
        for i in range(self.k):
            neighbours.append(distances[i][0])

        classes = {}
        for i in range(len(neighbours)):
            response = neighbours[i][-1]
            if response in classes:
                classes[response] += 1
            else:
                classes[response] = 1

        sorted_classes = sorted(classes.items(), key=lambda x: x[1], reverse=True)
        return sorted_classes[0][0]

    @staticmethod
    def evaluate(y_true, y_pred):
        n_correct = 0
        for act, pred in zip(y_true, y_pred):
            if act == pred:
                n_correct += 1
        return n_correct / len(y_true)
```

```
In [13]: knn=KNeighbours(k=3)
preds=[]
```

```
In [14]: train_set=pd.concat([pd.DataFrame(X_train),pd.DataFrame(Y_train)],axis=1)
test_set=pd.concat([pd.DataFrame(X_test),pd.DataFrame(Y_test)],axis=1)
train_set=train_set.astype(float).values.tolist()
test_set=test_set.astype(float).values.tolist()
```

```
In [15]: for row in test_set:
          predictors = row[:-1]
          pred=knn.predict(train_set,predictors)
          preds.append(pred)
```

```
In [16]: actual = np.array(test_set)[:,-1]
          print("Accuracy: {} %".format(knn.evaluate(actual, preds)*100))
```

Accuracy: 84.3 %

```
In [17]: print(classification_report(Y_test,preds))
```

	precision	recall	f1-score	support
0	0.83	0.72	0.77	727
1	0.85	0.91	0.88	1273
accuracy			0.84	2000
macro avg	0.84	0.82	0.82	2000
weighted avg	0.84	0.84	0.84	2000

## Inference

Classification Accuracy of K-Neighbours algorithm (Library) : 91.65 % with k=3

Classification Accuracy of K-Neighbours algorithm (Custom) : 84.30 % with k=3

**Inference** : The library function is better optimised in terms of the prediction subroutine than the custom written function as the custom written function.

# Clustering

## Dataset Description

### Anuran Calls (MFCCs) Data Set

Acoustic features extracted from syllables of anuran (frogs) calls, including the family, the genus, and the species labels (multilabel).

```
In [1]: import pandas as pd
import numpy as np
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
```

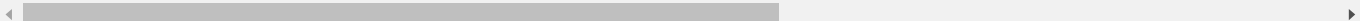
```
In [2]: data=pd.read_csv("Frogs_MFCCs.csv")
```

```
In [3]: data.head()
```

```
Out[3]:
```

	MFCCs_1	MFCCs_2	MFCCs_3	MFCCs_4	MFCCs_5	MFCCs_6	MFCCs_7	MFCCs_8	MFCCs_9	MFCCs_10	...	MFCCs_17	MFCCs_18
0	1.0	0.152936	-0.105586	0.200722	0.317201	0.260764	0.100945	-0.150063	-0.171128	0.124676	...	-0.108351	-0.077623
1	1.0	0.171534	-0.098975	0.268425	0.338672	0.268353	0.060835	-0.222475	-0.207693	0.170883	...	-0.090974	-0.056510
2	1.0	0.152317	-0.082973	0.287128	0.276014	0.189867	0.008714	-0.242234	-0.219153	0.232538	...	-0.050691	-0.023590
3	1.0	0.224392	0.118985	0.329432	0.372088	0.361005	0.015501	-0.194347	-0.098181	0.270375	...	-0.136009	-0.177037
4	1.0	0.087817	-0.068345	0.306967	0.330923	0.249144	0.006884	-0.265423	-0.172700	0.266434	...	-0.048885	-0.053074

5 rows × 26 columns



```
In [4]: X=data.iloc[:,1:-4]
Y=data.iloc[:, -4]
```

```
In [5]: Y.nunique()
```

```
Out[5]: 4
```

```
In [6]: sc=StandardScaler()
X=sc.fit_transform(X)
```

```
In [7]: X_train,X_test,Y_train,Y_test=train_test_split(X,Y,test_size=0.2,random_state=0)
```

## K-Means Clustering

### K-means Clustering (Library)

```
In [8]: from sklearn.cluster import KMeans
```

```
In [9]: from sklearn.metrics import silhouette_score

x=[]
sil = []
kmax = 10

for k in range(2, kmax+1):
    kmeans = KMeans(n_clusters = k).fit(X_train)
    labels = kmeans.labels_
    sil.append(silhouette_score(X_train, labels, metric = 'euclidean'))
    x.append(k)
```

```
In [10]: for i in zip(x,sil):  
         print(i)
```

```
(2, 0.33615042568302983)  
(3, 0.3586470357055605)  
(4, 0.36090644028848684)  
(5, 0.36469802624333525)  
(6, 0.2825279468442282)  
(7, 0.28995488336593667)  
(8, 0.2932980160754407)  
(9, 0.29943154667052946)  
(10, 0.23809084635969954)
```


```
In [11]: import matplotlib.pyplot as plt  
import seaborn as sns  
from sklearn.decomposition import PCA  
km = KMeans(n_clusters = 5)
```

```
In [12]: data=pd.DataFrame(X_train)  
data['clusters'] = km.fit_predict(data)  
reduced_data = PCA(2).fit_transform(data)
```

```
In [13]: reduced_data.shape
```

```
Out[13]: (5756, 2)
```

```
In [14]: results = pd.DataFrame(reduced_data,columns=['pca1','pca2'])  
sns.scatterplot(x="pca1", y="pca2", hue=data['clusters'], data=results)  
plt.title('K-means Clustering (Library)')  
plt.show()
```



```
In [15]: from sklearn.metrics import silhouette_score  
silhouette_score(X_train,data['clusters'])
```

```
Out[15]: 0.3647438484624485
```

## K-Means Clustering (Custom)

(continued on next page)

```
In [16]: class KMeans_custom:
def __init__(self, n_clusters):
    self.data = pd.DataFrame()
    self.n_clusters = n_clusters
    self.centroids = pd.DataFrame()
    self.clusters = np.ndarray(1)
    self.old_centroids = pd.DataFrame()
    self.verbose = False
    self.predictions = list()

def train(self, df, verbose):
    self.verbose = verbose
    self.data = df.copy(deep=True)
    self.clusters = np.zeros(len(self.data))

    if 'species' in self.data.columns:
        self.data.drop('species', axis=1, inplace=True)

    unique_rows = self.data.drop_duplicates()
    unique_rows.reset_index(drop=True, inplace=True)
    self.centroids = unique_rows.sample(n=self.n_clusters)
    self.centroids.reset_index(drop=True, inplace=True)

    if self.verbose:
        print("\nRandomly initiated centroids:")
        print(self.centroids)

    self.old_centroids = pd.DataFrame(np.zeros(shape=(self.n_clusters, self.data.shape[1])),
                                       columns=self.data.columns)

    while not self.old_centroids.equals(self.centroids):

        if self.verbose:
            time.sleep(3)

        self.old_centroids = self.centroids.copy(deep=True)

        for row_i in range(0, len(self.data)):
            distances = list()
            point = self.data.iloc[row_i]

            for row_c in range(0, len(self.centroids)):
                centroid = self.centroids.iloc[row_c]
                distances.append(np.linalg.norm(point - centroid))

            self.clusters[row_i] = np.argmin(distances)

        for cls in range(0, self.n_clusters):

            cls_idx = np.where(self.clusters == cls)[0]

            if len(cls_idx) == 0:
                self.centroids.loc[cls] = self.old_centroids.loc[cls]
            else:
                # Set the new k-mean to the mean value of the data points within this cluster
                self.centroids.loc[cls] = self.data.iloc[cls_idx].mean()

            if self.verbose:
                print("\nRow indices belonging to cluster {}: [n={}]".format(cls, len(cls_idx)))
                print(cls_idx)


        if self.verbose:
            print("\nOld centroids:")
            print(self.old_centroids)
            print("\nNew centroids:")
            print(self.centroids)
```

```
In [17]: km = KMeans_custom(n_clusters=5)
```

```
In [18]: data=pd.DataFrame(X_train)
km.train(data,verbose=False)
```

```
In [19]: data['clusters'] = km.clusters
reduced_data = PCA(2).fit_transform(data)
```

```
In [20]: results = pd.DataFrame(reduced_data, columns=['pca1', 'pca2'])
sns.scatterplot(x="pca1", y="pca2", hue=data['clusters'], data=results)
plt.title('K-means Clustering (custom)')
plt.show()
```



```
In [21]: silhouette_score(X_train, data['clusters'])
```

```
Out[21]: 0.19872144546341883
```

## Inference

K-Means Clustering (Library) Silhouette Score : 0.36 where n=5

K-Means Clustering (Custom) Silhouette Score : 0.19 where n=5

Lower silhouette score of the custom algorithm indicates more overlapping of the clusters, thereby the library function gets a better score due to it's optimised techniques.

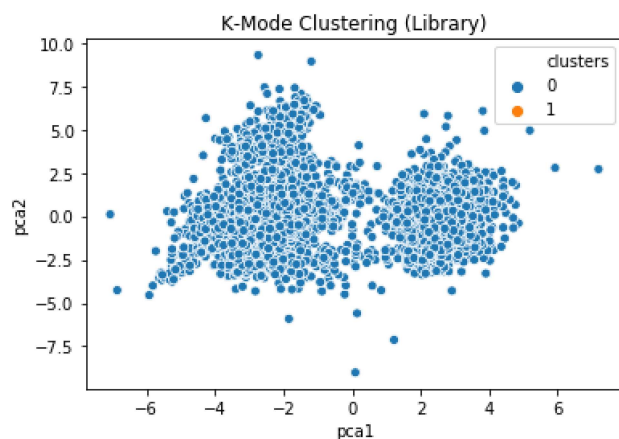
## K-Mode Clustering

### K-Mode Clustering (Library)

```
In [22]: from kmodes.kmodes import KModes
kmo = KModes(n_clusters = 2)
```

```
In [23]: data=pd.DataFrame(X_train)
data['clusters'] = kmo.fit_predict(data)
reduced_data = PCA(2).fit_transform(data)
```

```
In [24]: results = pd.DataFrame(reduced_data, columns=['pca1', 'pca2'])
sns.scatterplot(x="pca1", y="pca2", hue=data['clusters'], data=results)
plt.title('K-Mode Clustering (Library)')
plt.show()
```



```
In [25]: silhouette_score(X_train,data['clusters'])
```

```
Out[25]: 0.18575828514490955
```

## K-Mode Clustering (Custom)

```
In [26]: from sklearn.base import BaseEstimator, ClusterMixin
from sklearn.utils import check_random_state
from sklearn.utils.validation import check_array
```

```
In [27]: from util import get_max_value_key, encode_features, get_unique_rows, decode_centroids, pandas_to_numpy
from util.dissim import matching_dissim, ng_dissim
```

```
In [28]: from helper import init_huang,init_cao,move_point_cat,_labels_cost,_k_modes_iter,k_modes,k_modes_single
```

```
In [29]: class KModes_Custom(BaseEstimator, ClusterMixin):
    def __init__(self, n_clusters=8, max_iter=100, cat_dissim=matching_dissim,
                 init='Cao', n_init=1, verbose=0, random_state=None, n_jobs=1):

        self.n_clusters = n_clusters
        self.max_iter = max_iter
        self.cat_dissim = cat_dissim
        self.init = init
        self.n_init = n_init
        self.verbose = verbose
        self.random_state = random_state
        self.n_jobs = n_jobs
        if ((isinstance(self.init, str) and self.init == 'Cao') or
            hasattr(self.init, '__array__')) and self.n_init > 1:
            if self.verbose:
                print("Initialization method and algorithm are deterministic. "
                      "Setting n_init to 1.")
            self.n_init = 1

    def fit(self, X, y=None, **kwargs):
        X = pandas_to_numpy(X)

        random_state = check_random_state(self.random_state)
        self._enc_cluster_centroids, self._enc_map, self.labels_, self.cost_, \
        self.n_iter_, self.epoch_costs_ = k_modes(
            X,
            self.n_clusters,
            self.max_iter,
            self.cat_dissim,
            self.init,
            self.n_init,
            self.verbose,
            random_state,
            self.n_jobs,
        )
        return self

    def fit_predict(self, X, y=None, **kwargs):
        return self.fit(X, **kwargs).predict(X, **kwargs)

    def predict(self, X, **kwargs):
        assert hasattr(self, '_enc_cluster_centroids'), "Model not yet fitted."
        X = pandas_to_numpy(X)
        X = check_array(X, dtype=None)
        X, _ = encode_features(X, enc_map=self._enc_map)
        return _labels_cost(X, self._enc_cluster_centroids, self.cat_dissim)[0]

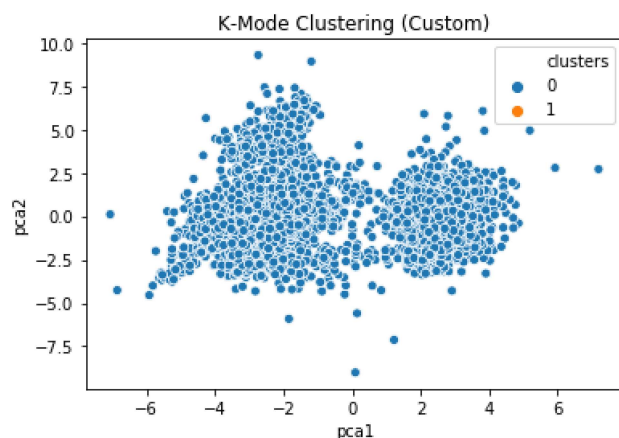
    @property
    def cluster_centroids_(self):
        if hasattr(self, '_enc_cluster_centroids'):
            return decode_centroids(self._enc_cluster_centroids, self._enc_map)
        else:
            raise AttributeError("'{}' object has no attribute 'cluster_centroids_' "
                                  "because the model is not yet fitted.")
```

```
In [30]: kmo = KModes_Custom(n_clusters = 2)
```

```
In [31]: data=pd.DataFrame(X_train)
data['clusters'] = kmo.fit_predict(data)
reduced_data = PCA(2).fit_transform(data)
```



```
In [32]: results = pd.DataFrame(reduced_data, columns=['pca1', 'pca2'])
sns.scatterplot(x="pca1", y="pca2", hue=data['clusters'], data=results)
plt.title('K-Mode Clustering (Custom)')
plt.show()
```



```
In [34]: silhouette_score(X_train, data['clusters'])
```

Out[34]: 0.11330828514490955

## Inference

K-Modes Clustering (Library) Silhouette Score : 0.18 where n=2

K-Modes Clustering (Custom) Silhouette Score : 0.11 where n=2

- The low performance of k-mode clustering algorithms is due to the fact that k-mode is more suited for categorical variables and this dataset lacks categorical variables.
- Lower silhouette score of the custom algorithm indicates more overlapping of the clusters, thereby the library function gets a better score due to its optimised techniques.